

Using PAPI to Monitor Communication Between Processes Running in Parallel

Daniel Topa

April 25, 2025

Abstract

This document demonstrates how to use the **Performance Application Programming Interface (PAPI)** to profile inter-process communication in parallel applications. We cover integration with **MPI**, **Coarray Fortran**, and **Legion**, compare PAPI with alternative tools, and provide ready-to-use code examples. Includes installation via Spack and advanced profiling toolchains.

1 Installation

1.1 Installing PAPI via Spack

```
1 spack install papi
```

Listing 1: Installing PAPI with Spack

2 Monitoring MPI Communication

```
1 #include <mpi.h>
2 #include <papi.h>
3
4 int main(int argc, char **argv) {
5     MPI_Init(&argc, &argv);
6     int rank;
7     MPI_Comm_rank(MPI_COMM_WORLD, &rank);
8
9     // Initialize PAPI
10    int events[2] = {PAPI_TOT_CYC, PAPI_L2_TCM};
11    long long values[2];
12    PAPI_start_counters(events, 2);
13
14    // Communication pattern
15    double data[1000];
16    if (rank == 0) {
17        MPI_Send(data, 1000, MPI_DOUBLE, 1, 0, MPI_COMM_WORLD);
18    } else {
19        MPI_Recv(data, 1000, MPI_DOUBLE, 0, 0, MPI_COMM_WORLD,
20                MPI_STATUS_IGNORE);
21    }
22
23    // Read counters
24    PAPI_stop_counters(values, 2);
25    printf("Rank %d: Cycles=%lld, L2 Misses=%lld\n",
26           rank, values[0], values[1]);
27}
```

```

28     MPI_Finalize();
29     return 0;
30 }

```

Listing 2: MPI+PAPI Example

2.1 Compilation and Execution

```

1 mpicc -o mpi_papi mpi_papi.c -lpapi
2 mpiexec -n 2 ./mpi_papi

```

2.2 Expected Output

Rank 0: Cycles=1200000, L2 Misses=850
Rank 1: Cycles=1800000, L2 Misses=920

3 Monitoring Coarray Fortran

3.1 Example Code

```

1 program caf_papi
2   use iso_c_binding
3   implicit none
4   integer(c_int) :: me, np
5   real :: array[*]
6
7   interface
8     subroutine papi_start() bind(C)
9     end subroutine
10    function papi_read_cycles() bind(C) result(cycles)
11      import :: c_long_long
12      integer(c_long_long) :: cycles
13    end function
14  end interface
15
16  me = this_image()
17  np = num_images()
18
19  call papi_start()
20  array = me * 2.0
21  sync all ! Measure synchronization cost
22  print *, "Image", me, "Cycles:", papi_read_cycles()
23 end program

```

Listing 3: Coarray+PAPI Example

3.2 Compilation and Execution

```

1 caf -lpapi caf_papi.f90 -o caf_papi
2 cafrun ./caf_papi # assumes chmod +x caf_papi

```

4 Complementary Profiling Tools

4.1 Tracing Tools

- **VampirTrace/Score-P:** Combines MPI, OpenMP, and Coarray Fortran tracing

```

1 spack install scorep
2 export SCOREP_ENABLE_TRACING=1
3 scorep --caf mpif90 -o caf_prog caf_prog.f90

```

- **TAU (Tuning and Analysis Utilities):**

```

1 spack install tau
2 tau_f90.sh -o caf_tau caf_prog.f90

```

4.2 Call-Path Profiling

- **HPCToolkit:**

```

1 spack install hpctoolkit
2 hpcrun -e PAPI_TOT_CYC ./caf_prog
3 hpcprof -S caf_prog.hpcstruct hpctoolkit-caf_prog-*

```

- **Legion Prof:**

- Built into Legion runtime (enable with `-lg:prof`)
- Visualizes task graphs and dependencies
- Compatible with PAPI counters via `-lg:papi`

5 Tool Comparison

Tool	Strengths	Integration
PAPI	Hardware counters	Manual instrumentation
Score-P	Full MPI/CAF tracing	Automated via compiler
TAU	Portable profiling	Runtime configuration
HPCToolkit	Call-path analysis	Post-processing
Legion Prof	Task visualization	Native to Legion

Table 1: Advanced Profiling Tools Comparison

```

1 # Run Legion application with PAPI
2 LEGION_PAPI_COUNTERS=PAPI_TOT_CYC,PAPI_FP_OPS ./legion_prog -lg:papi 1

```

Listing 4: Legion with PAPI Profiling

Appendix A: Spack Package Information

PAPI Package Details

The following shows the output of `spack info papi`, which provides comprehensive information about available versions, variants, and dependencies:

```

==> Warning: The packages:all:compiler preference has been deprecated...
AutotoolsPackage:  papi

Description:
  PAPI provides the tool designer and application engineer with a
  consistent interface and methodology for use of the performance counter
  hardware found in most major microprocessors...

```

Homepage: <https://icl.utk.edu/papi/>

Preferred version:

7.1.0 <https://icl.utk.edu/projects/papi/downloads/papi-7.1.0.tar.gz>

Safe versions:

master [git] <https://github.com/icl-utk-edu/papi> on branch master
7.1.0 <https://icl.utk.edu/projects/papi/downloads/papi-7.1.0.tar.gz>
7.0.1 <https://icl.utk.edu/projects/papi/downloads/papi-7.0.1.tar.gz>
6.0.0.1 <https://icl.utk.edu/projects/papi/downloads/papi-6.0.0.1.tar.gz>
[additional versions...]

Variants:

build_system [autotools]	autotools
cuda [false]	false, true
debug [false]	false, true
example [true]	false, true
infiniband [false]	false, true
lmsensors [false]	false, true
nvml [false]	false, true
powercap [false]	false, true
rapl [false]	false, true
rocm [false]	false, true
rocm_smi [false]	false, true
sde [false]	false, true
shared [true]	false, true
static_tools [false]	false, true
when +rocm	
amdgpu_target [none]	none, gfx1010, gfx1011, [additional targets...]
when @6.0.0:	
rdpmc [true]	false, true

Build Dependencies:

bc c cuda cxx fortran gmake gnuconfig hip hsa-rocr-dev llvm-amdgpu
lm-sensors rocm-openmp-extras rocm-smi-lib rocprofiler-dev

Link Dependencies:

cuda hip hsa-rocr-dev llvm-amdgpu lm-sensors rocm-openmp-extras
rocm-smi-lib rocprofiler-dev

Run Dependencies:

None

Licenses:

BSD-3-Clause

Key observations:

- Latest stable version: 7.1.0
- Important variants: debug, cuda, rocm for GPU support
- ROCm requires specific AMD GPU architecture specification
- BSD-3-Clause license allows flexible use